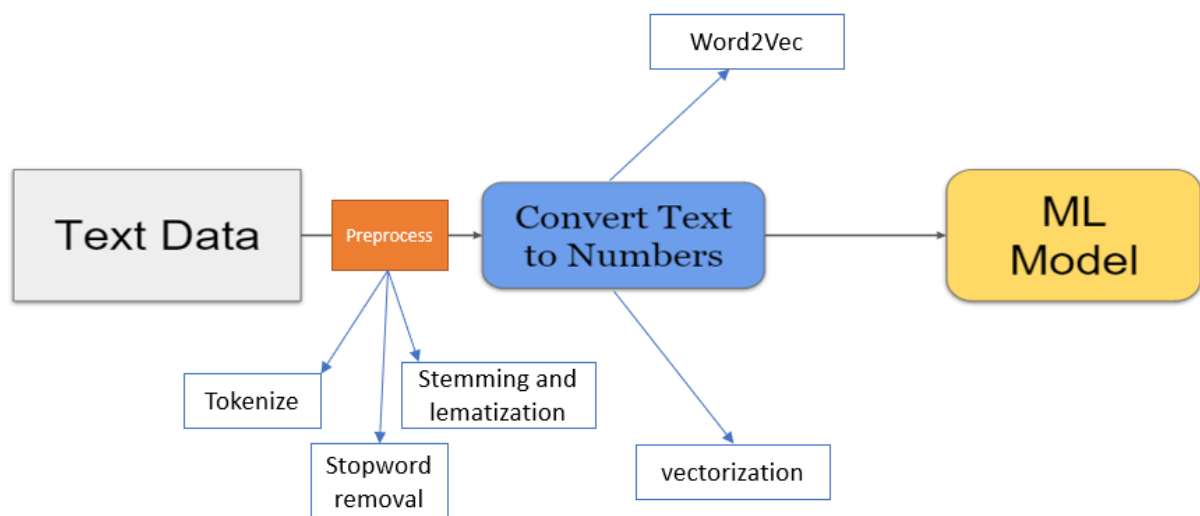


# NLP Rough Notes

- Corpus: A body of text samples
- Document: A text sample
- Vocabulary: A list of words used in the corpus
- Language model: How the words are supposed to be organized



## Tokenization

Chopping up text into pieces called tokens. Each word in a sentence is token.

### Method 1

```
docs = tweets['text'].str.lower().str.replace('[^a-z\s#@]', '') # remove every
thing other than alphabets, spaces, # , @
docs_tokens = docs.str.split(' ')

tokens_all = []
for tokens in docs_tokens:
    tokens_all.extend(tokens)
print('No. of tokens in entire corpus:', len(tokens_all))
```

### Method 2

```
from keras.preprocessing.text import Tokenizer

tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_x)
```

## Stop words

- Articles: a, an, the
- Common verbs: is, was, are
- Pronouns: he, she, it
- Conjunctions: for, and
- Prepositions: at, on, with

### Method 1

```
import nltk # natural language tool kit
nltk.download('stopwords')
common_stopwords = nltk.corpus.stopwords.words('english')
custom_stopwords = ['amp', 'rt']
all_stopwords = np.hstack([common_stopwords, custom_stopwords])
df_tokens = pd.DataFrame(tokens_freq).reset_index().rename(columns={'index': 'token', 0: 'frequency'})
df_tokens = df_tokens[~df_tokens['token'].isin(all_stopwords)]
```

### Method 2

```
from gensim.parsing.preprocessing import remove_stopwords

remove_stopwords('this movie is really pathetic')
```

```
'movie pathetic'
```

## Word Cloud

```
docs = tweets['text']
docs_strings = ' '.join(docs)
wc = WordCloud(background_color='white', stopwords=all_stopwords).generate(docs_strings)
plt.figure(figsize=(20,5))
plt.imshow(wc)
plt.axis('off');
```

## Stemming

Stemming –chopping off the end of words

Nannies become nanni

Caresses become caress

## PorterStemmer

```
from gensim.parsing.preprocessing import PorterStemmer
docs = imdb['review'].str.lower().str.replace('[^a-z\s]', '')
docs = docs.apply(remove_stopwords)
docs = stemmer.stem_documents(docs)
```

## Lemmatization

Lemma of a word is a more exact task than stemming. Perform lemma rather than stemming.

### Method 1


```
import spacy
nlp = spacy.load("en_core_web_sm")
doc = imdb['review'].iloc[0]

proc_doc = nlp(doc)
for token in proc_doc:
    print(token, '|', token.lemma_, '|', token.pos_)
```

### Method 2

```
from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()
lemmatizer.lemmatize("rocks")
```

## Vectorization



Document #1

He is a good boy. She is also good.

Document #2

Radhika is a good person.

Vocabulary

a, also, boy, good, He, is, person, She, Radhika

	a	also	boy	good	He	is	person	She	Radhika
Index	0	1	2	3	4	5	6	7	8
Document #1	1	1	1	2	1	2	0	1	0
<b>Document #2</b>	1	0	0	1	0	1	1	0	1

## CountVectorizer

```
from sklearn.feature_extraction.text import CountVectorizer
docs = imdb['review'].str.lower().str.replace('[^a-z\s]', '')
train_docs, test_docs = train_test_split(docs, test_size=0.2, random_state=1)
stopwords = nltk.corpus.stopwords.words('english')
vectorizer = CountVectorizer(stop_words=stopwords, min_df=10).fit(train_docs)
train_dtm = vectorizer.transform(train_docs)
test_dtm = vectorizer.transform(test_docs)
```

## TfidfVectorizer

Advantage: Gives less weightage to most occurring words in the corpus. Why?

Frequently occurring words are not significant for differentiating docs.

*For an example and formula check GL 'Session 1.pdf' page 50.*

```
from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer = TfidfVectorizer(min_df=5).fit(train_docs)
train_dtm_tfidf = vectorizer.transform(train_docs)
test_dtm_tfidf = vectorizer.transform(test_docs)
```

# Sentiment analysis

## ML Model

```
naive_bayes_model = MultinomialNB().fit(train_dtm, train_y)
test_y_pred = naive_bayes_model.predict(test_dtm)
print('Accuracy score: ', accuracy_score(test_y, test_y_pred))
print('F1 score: ', f1_score(test_y, test_y_pred, pos_label='negative'))
```

## Rule Based Algorithm

```
from nltk.sentiment.vader import SentimentIntensityAnalyzer
analyzer = SentimentIntensityAnalyzer()
```

```
review = 'i hate tea and i love cofee'
analyzer.polarity_scores(review)
```

```
{'neg': 0.339, 'neu': 0.275, 'pos': 0.385, 'compound': 0.128}
```

## Calculation of neg, pos, neu and compound scores

Take the above sentence as an example 'I hate tea and I love coffee'.

### Step 1:

Each word is given a score/weight according to VEDAR sentiment. Check the full list [here](#).

I -> Ignored (stopword)

Hate -> -2.7

Tea, and, coffee -> 0

Love -> 3.2

### Step 2:

Increment weights by 1

I -> Ignored (stopword)

Hate -> -3.7

Tea, and, coffee -> 1

Love -> 4.2

### Step 3:

Total = 3.7 + 1 + 1 + 1 + 4.2 = 10.9

Pos = % of positive score =  $4.2/10.9 = 0.385$

Neg = % of negative score =  $3.7/10.9 = 0.339$

Neu = % of neutral score =  $3/10.9 = 0.275$

compound scores =

$$\frac{\text{Total score}}{\sqrt{\text{Total score}^2 + \underset{\substack{\uparrow \\ 15}}{\text{alpha}}}}$$

= total score before step 2 (increment) is 0.5

=  $0.5 / \text{np.sqrt}(\text{np.square}(0.5) + 15)$

= 0.128

# Word2Vec

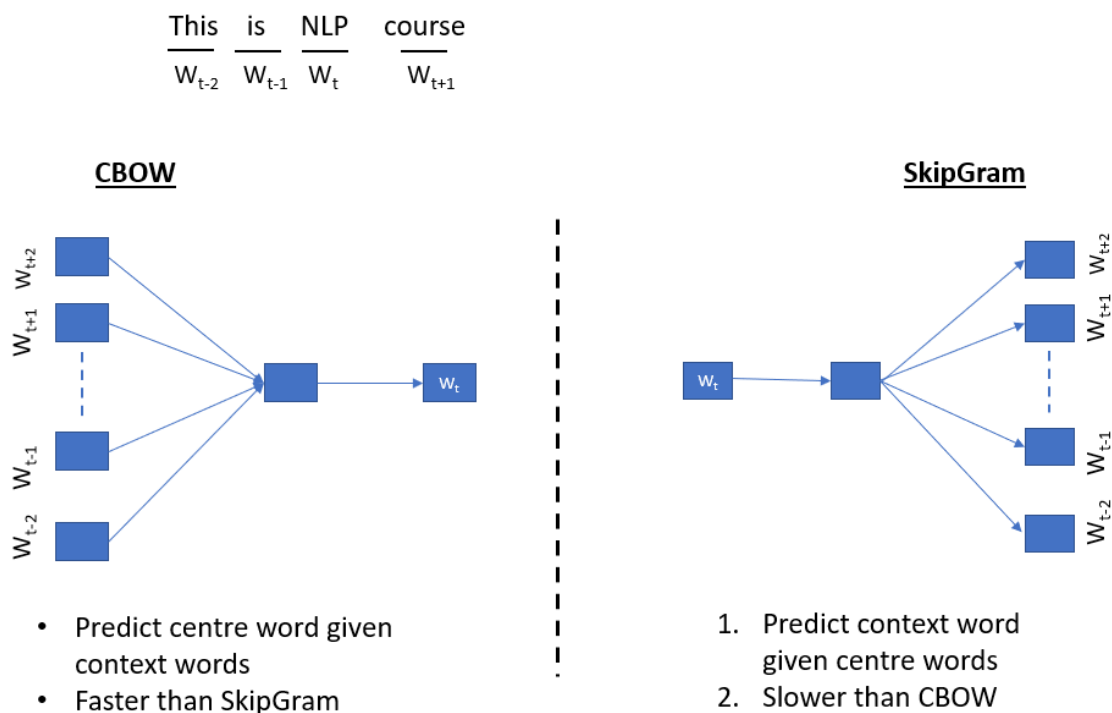
Overcome disadvantages of vectorization:

1. The order of the words in a sentence is not considered
2. Context of a sentence gets missed out

Objective:

1. Should be dense, not sparse like vectorization methods
2. Lower dimension, typically 300 vocabs
3. Represent meaning of the word
4. Should be comparable with each other

Two methods: **Count Bag of Words** and SkipGram



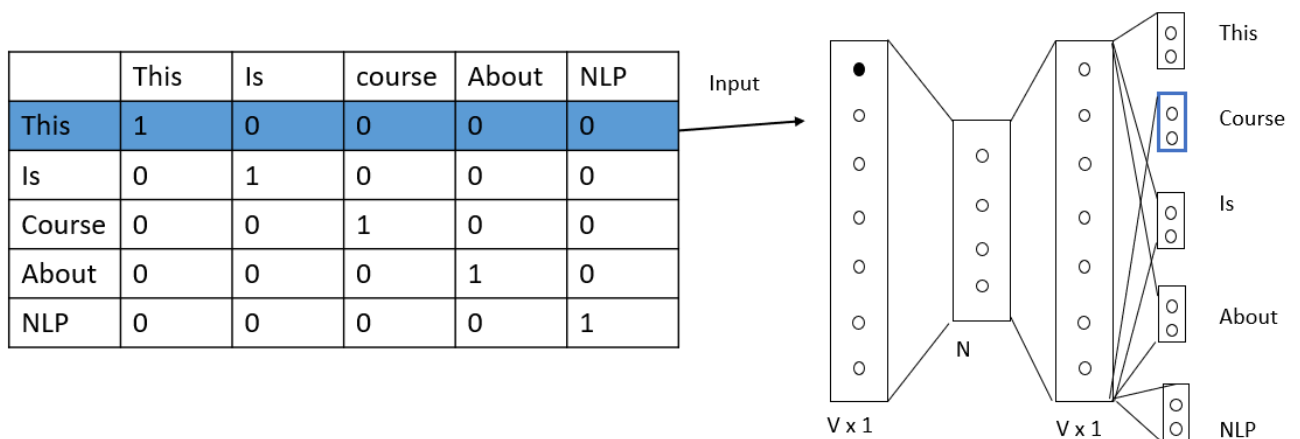
## SkipGram

Sentence: 

This	Course	is	about	NLP
------	--------	----	-------	-----

With sliding window = 1


Centre word	Context words
This	Course
Course	(this, is)
Is	(course, about)
About	(is, <u>nlp</u> )
NLP	about



V -> vocab size

N -> hidden layer nodes

Row-wise input to the neural network. When row 1 is fed, 'course' dense layer is activated. Above this is initial architecture. Below is the current architecture used.



- Output will be word embedding matrix.
- This method is computationally challenging due to large dense layers at output end.

## Negative sampling to overcome the challenge

Centre word	Context words		Centre word	Context words	
This	Course	Add noise 1 context pair to 1 context word	This	Course	+ve
Course	(this, is)		this	math	-ve/noise
Is	(course, about)		course	this	+ve
About	(is, nlp)		Course	is	-ve/noise
NLP	about		course	science	-ve/noise

- Convert context pair into single context word
- Introduce noise, meaning, context word that does not exist in the document and label it as -ve or 0.
- Context word that appears in the doc will be labeled +ve or 1
- Now we just have to build a binary classifier neural network that predicts 1 or 0. This significantly reduced complexity at the output layer.
- A lot of pre-trained models are available to use.

## Pre-trained Glove Word2Vec embedding

```
glove_path = 'glove.840B.300d/glove.840B.300d.txt'

with zf.open(glove_path) as file:
    embeddings = {}
    for line in file:
        line = line.decode('utf-8').replace('\n', '').split(' ')
        word = line[0]
        if word in vocab:
            vector = line[1:]
            vector = [float(x) for x in vector]
            embeddings[word] = vector
    embedding_dim = len(vector)
    embedding_matrix = np.zeros((vocab_size, embedding_dim))
    for word, index in tokenizer.word_index.items():
        if word in embeddings:
            embedding_matrix[index] = embeddings[word]

model = Sequential()
model.add(Embedding(vocab_size, embedding_dim, weights=[embedding_matrix]))
.
```

## CBOW

Exactly opposite of SkipGram



## Self-trained SkipGram and CBOW example

# Create CBOW model

```
model1 = gensim.models.Word2Vec(data, min_count = 1, size = 100, window = 5)
```


```
model1.similarity('alice', 'wonderland')
```

# Create Skip Gram model

```
model2 = gensim.models.Word2Vec(data, min_count = 1, size = 100, window = 5, sg = 1)
```


```
model2.similarity('alice', 'wonderland')
```

## RNN



- RNN is a compressed version of ANN.
- It has a temporal loop in between.
- Common way to represent is to unwind the loop (last fig).
- This architecture will allow the network to have memory of previous input.

Disadvantages:




1. Vanishing gradient
2. **Long-Term Dependency Problem:** Information of  $x_0$  will get lost by the time it reaches output node
3. Middle data ( $x_1$  or  $x_3$  or ...) might not be helpful

LSTM overcomes these disadvantages.

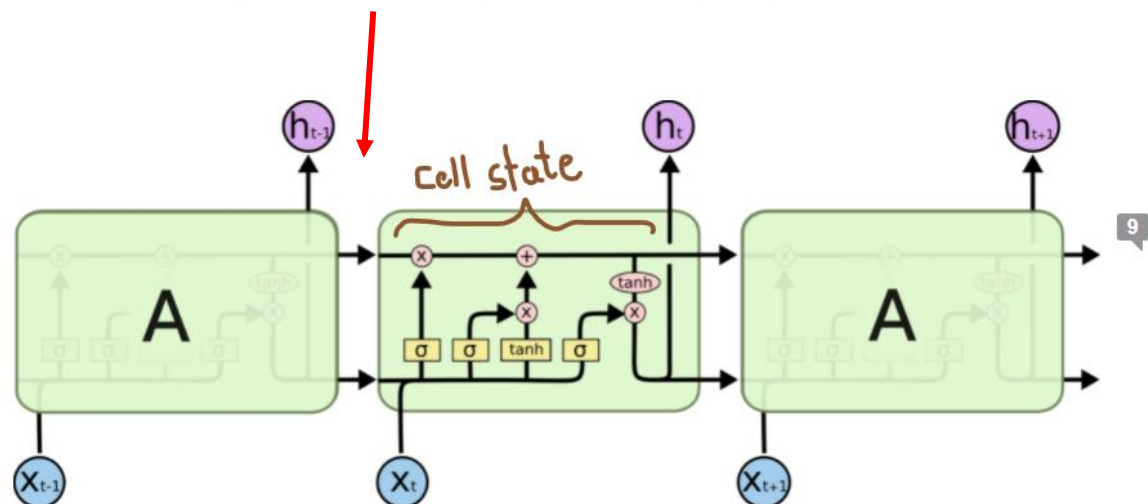
## LSTM

LSTMs are explicitly designed to **avoid the long-term dependency problem**. Remembering information for long periods of time is practically their default behavior, not something they struggle to learn!

- They are carefully regulated by structures called **gates** (yellow boxes in below diagram)
- Gates are a way to optionally let information through. They are composed out of a sigmoid neural net layer and a pointwise multiplication operation.
- The sigmoid layer outputs numbers between zero and one, describing how much of each component should be let through. A value of zero means “let nothing through,” while a value of one means “let everything through!”
- The key to LSTMs is the **cell state**, the horizontal line running through the top of the diagram.




The repeating module in a standard RNN contains a single layer.




The repeating module in an LSTM contains four interacting layers.

Notations:




Steps 1:




- The first step in our LSTM is to decide what information we're going to throw away from the cell state.
- This decision is made by a sigmoid layer called the “**forget gate layer**.” It looks at  $h_{t-1}$  and  $x_t$ , and outputs a number between 0 and 1 for each number in the cell state  $C_{t-1}$ .

### Step 2:




- The next step is to decide what new information we're going to store in the cell state.
- First, a sigmoid layer called the “**input gate layer**” decides which values we'll update.
- tanh layer creates a vector of new candidate values,  $\tilde{C}_t$

### Step 3:



- Update the old cell state,  $C_{t-1}$ , into the new cell state  $C_t$ .
- We multiply the old state by  $f_t$ , forgetting the things we decided to forget earlier.
- Then we add  $i_t \cdot \tilde{C}_t$ . This is the new candidate values, scaled by how much we decided to update each state value.

### Step 4:



- Decide what we're going to output
- sigmoid layer which decides what parts of the cell state we're going to output.
- Then, we put the cell state through tanh (to push the values to be between  $-1$  and  $1$ ) and multiply it by the output of the sigmoid gate, so that we only output the parts we

# units: no.of LSTM memory cells(neurons)

# return\_sequences: True if more than 1 LSTM layers

```
model.add(LSTM(units = 50, return_sequences = True, input_shape = (X_train.shape[1], 1)))
```

```
model.add(LSTM(units = 50, return_sequences = False))
```


#output layer

```
model.add(Dense(units = 1))
```

- Opening and closing of gates is not controlled manually. The neural network will become smart enough to control it over time.
- Hence, each gate is a separate neural network.
- Each input( $X_{t-1}$ ,  $X_t$ ,  $X_{t+1}$ ) and output( $h_{t-1}$ ,  $h_t$ ,  $h_{t+1}$ ) nodes are not just a single node. They are series of nodes, one behind the other, not visible in 2D diagrams.

## GRU


LSTM has many gates, which results in overfitting most of the time. GRU is a simplified version of LSTM with a fewer number of gates.



Amazing intuition video of RNN, LSTM and GRU. Animated explanation of working of gates: <https://www.youtube.com/watch?v=8HyCNIVRbSU>

# Attention method

Alternate structure of RNN. What attention is given to each input and process the model in parallel. This attention method will be used in Encoder-Decoder.




- Sequential operation. Without computing  $h_1$  you can not compute  $h_2$ . Interdependent
- Hence, time to compute RNN is very high

- Separate parallel units for each  $h_i$ .
- Every input is subsequently passed.
- At every instance  $h_i$ , we add new input  $x_i$ .
- Execute each unit in parallel.
- Scoring mechanism to each input  $x_i$ . Weights of each input  $x_i$  is computed by the model.
- The weight represents how much of information from each  $x_i$  is passed to the network.


## Encoder-Decoder

Another variant of RNN. Vector to a vector model. Applications:

1. Machine translation (one language to another)
2. sequence-to-sequence modeling
3. Natural language query -> SQL query
4. Natural language -> program

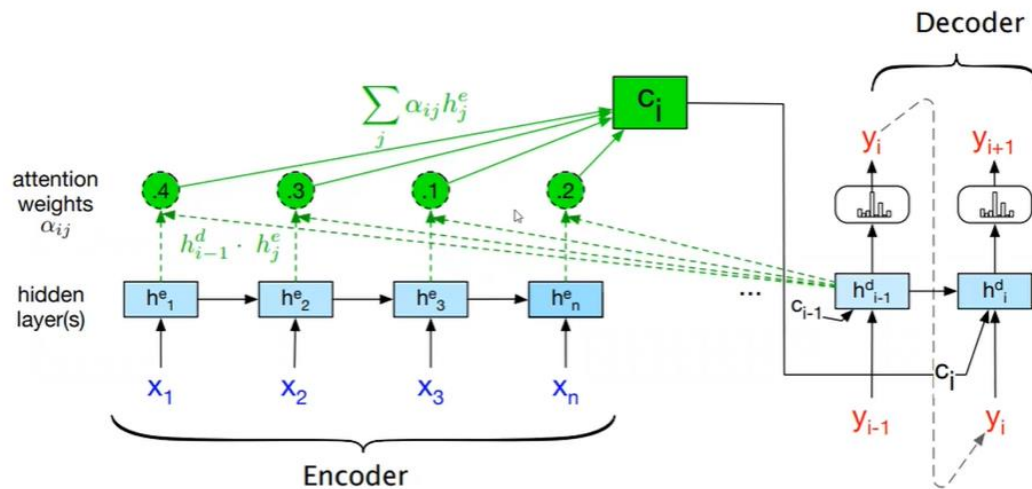


- Consider an example for translating english sentence "I like tea" into German words " $G_1 G_2 G_3 G_4$ ".
- Notice, number of words in output sentence might to be equal to number of words in input sentence.
- Which is why, words by words prediction is not performed, that is,  $I \rightarrow G_1$ , like  $\rightarrow G_2$ . Instead process the entire sentence and then start prediction.
- What we get at the end of processing entire sentence is called a context.
- This approach has two RNN's one for encoder and one for decoder.



- Encoder is to process input sentence and output a context.
- Think of encoder like a PCA, consolidate all the feature values and the context vector as different PC's
- $H_0$  to encoder is randomly initializer or set to zero.
- Input for decoder will be the context which is used to predict the German words.
- Each predicted German word is inputed to the next unit in the network.
- Decoder is where we have attention mechanism.

## Attention mechanism in the decoder



- We learned that we need the final output context as input for the decoder.
- In the improved version, we should also pass each intermediate hidden state of the encoder as decoder input.
- There intermediate hidden state should be weighted. This is where the attention mechanism kicks in.

Encoder-Decoder code with comments [link](#)